

HW 5

Q1

1.1

```
#rm(list=ls())
csv1 <- read.csv('model1.csv')
summary(lm(y ~ d + x1, data = csv1))

##
## Call:
## lm(formula = y ~ d + x1, data = csv1)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.6680 -0.2548 -0.2548  0.3320  0.9564
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.043571   0.005051   8.626  <2e-16 ***
## d            0.413189   0.005266  78.466  <2e-16 ***
## x1           0.211207   0.005127  41.191  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.4321 on 99997 degrees of freedom
## Multiple R-squared:  0.05828,    Adjusted R-squared:  0.05826
## F-statistic: 3094 on 2 and 99997 DF,  p-value: < 2.2e-16
```

ATE is 0.413. We want to adjust for non-causal correlation, X1, but not block causal paths through X2.

1.2

```
csv2 <- read.csv('model2.csv')
summary(lm(y ~ d + x1 + x2, data = csv2))

##
## Call:
## lm(formula = y ~ d + x1 + x2, data = csv2)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.7527 -0.2126 -0.1937  0.2473  1.0669
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -0.066891   0.001959  -34.14  <2e-16 ***
## d            0.260606   0.003045   85.60  <2e-16 ***
```

```
## x1          0.279484    0.002508   111.45    <2e-16 ***
## x2          0.279500    0.002518   111.02    <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.3536 on 99996 degrees of freedom
## Multiple R-squared:  0.3915, Adjusted R-squared:  0.3914
## F-statistic: 2.144e+04 on 3 and 99996 DF,  p-value: < 2.2e-16
```

ATE is 0.260606. We want to adjust for non-causal correlation.

1.3

```
csv3 <- read.csv('model3.csv')
summary(lm(y ~ d + x1, data = csv3))

##
## Call:
## lm(formula = y ~ d + x1, data = csv3)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -0.7057 -0.5022  0.2943  0.3561  0.5596
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  0.440410    0.004020   109.56  <2e-16 ***
## d            0.061789    0.003266    18.92  <2e-16 ***
## x1           0.203493    0.003533    57.60  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.476 on 99997 degrees of freedom
## Multiple R-squared:  0.03213,    Adjusted R-squared:  0.03211
## F-statistic: 1660 on 2 and 99997 DF,  p-value: < 2.2e-16
```

ATE is 0.061789. We want to adjust for unknown confounders.

Q2

2.1

We do not know U – an unknown confounder. So, we cannot identify ATE via simple regression adjustment.

2.2

$Y_i = \lambda_{dy}D_i + U_i = \lambda_{dy}\lambda_{xd}U_{xi} + \lambda_{dy}U_i + U_i$. The bias is obviously larger when we have X adjusted: for $Y_i = \lambda_{dy}D_i + kX_i + U_i = \lambda_{dy}\lambda_{xd}U_{xi} + \lambda_{dy}U_i + U_i + k\lambda_{xd}U_{xi} + kU_i$. Therefore, there is a bias amplification effect.

2.3

$$Cov(X, Y) = Cov(X, \lambda_{dy}D_i + U_i) = Cov(X, \lambda_{dy}D_i) = \lambda_{dy}Cov(X, D_i)$$

$$\text{Thus, } \lambda_{dy} = \frac{Cov(X, Y)}{Cov(X, D_i)} = \frac{\frac{Cov(X, Y)}{Var(X)}}{\frac{Cov(X, D_i)}{Var(X)}} = \frac{\beta_{yx}}{\beta_{dx}}.$$

Q3

3.1

```
# function to compute k-fold CV MSE for polynomial regression
cv_mse <- function(data, d = 1, k = 10){

  # create folds randomly
  n <- nrow(data)
  folds <- sample(rep(1:k, length = n))

  # create vector to store results
  mse <- rep(NA, k)
  for(j in 1:k){

    # train model on all folds except j
    train <- folds != j
    ols <- lm(y ~ poly(x, degree = d, raw = T), data = data[train, ])

    # compute MSE on fold j (not used for training)
    yhat <- predict(ols, newdata = data[!train, ])
    mse[j] <- mean((data$y[!train] - yhat)^2)
  }

  # compute average mse
  mse.cv <- mean(mse)
  return(mse.cv)
}
```

Function

```
set.seed(42)
n = 100
x = runif(n, -4, 4)
e = rnorm(n)

##?ifelse

# simulate DGP1
y <- -2*ifelse(x < -3, 1,0) + 2.55 * ifelse(x > -2, 1,0) - 2*ifelse(x > 0, 1,0) +
  4 * ifelse(x > 2, 1,0) - ifelse(x > 3, 1,0) + e
sim_data1 <- data.frame(x, y)

# simulate DGP2
y <- 6 + 0.4 * x - 0.36 * x^2 + 0.006 * x^3 + e
sim_data2 <- data.frame(x, y)


# simulate DGP3
y <- 2.83*sin(pi/2*x) + e
sim_data3 <- data.frame(x, y)

# simulate DGP4
y <- 4*sin(3*pi *x) * ifelse(x > 0, 1,0) + e
sim_data4 <- data.frame(x, y)
```

DGP

```
# plot size
options(repr.plot.width = 10, repr.plot.height = 10)

# compute MSE's for d from 1 to 10
degree <- 1:10
cv.mse1 <- sapply(degree, function(d) cv_mse(sim_data1, d))
plot(degree, cv.mse1)
```



Sim data 1

```
# best degree
best1 <- degree[which.min(cv.mse1)]
best1
```

```
## [1] 8
```

```
# fit model using best degree
ols1 <- lm(y ~ poly(x, degree = best1, raw = T), data = sim_data1)
```

```
# predicted values
yhat1 <- predict(ols1)
```

```
# plot against data and true CEF
```


```
plot.new()
plot(y ~ x, sim_data1, pch= 20)
curve(-2*ifelse(x < -3, 1, 0) + 2.55 * ifelse(x > -2, 1, 0) - 2*ifelse(x > 0, 1, 0)
      + 4 * ifelse(x > 2, 1, 0) - ifelse(x > 3, 1, 0) + e, col = "red", from = -4, to = 4, add = T, lwd = 4)
```

```
## Warning in -2 * ifelse(x < -3, 1, 0) + 2.55 * ifelse(x > -2, 1, 0) - 2 * :
```

```
## longer object length is not a multiple of shorter object length
```

```
lines(yhat1[order(x)] ~ sort(x), data= sim_data1, col = "blue", lwd = 4)
```

```
legend("topleft", col = c("red", "blue"), legend = c("true", "fitted"), lty = 1, bty = "n", lwd = 4)
```




```
# compute MSE's for d from 1 to 10
degree <- 1:10
cv.mse2 <- sapply(degree, function(d) cv_mse(sim_data2, d))

# best degree
best2 <- degree[which.min(cv.mse2)]
best2
```

Sim data 2

```
## [1] 6
```

```
plot(degree, cv.mse2)
```




```
# fit model using best degree
ols2 <- lm(y ~ poly(x, degree = best2, raw = T), data = sim_data2)

# predicted values
yhat2 <- predict(ols2)

# plot against data and true CEF
plot.new()
plot(y ~ x, sim_data2, pch= 20)
curve( 6 + 0.4 * x - 0.36 * x^2 + 0.006 * x^3 + e, col = "red", from = -4, to = 4, add = T, lwd = 4)

## Warning in 6 + 0.4 * x - 0.36 * x^2 + 0.006 * x^3 + e: longer object length is
## not a multiple of shorter object length

lines(yhat2[order(x)] ~ sort(x), data= sim_data2, col = "blue", lwd = 4)
legend("topleft", col = c("red", "blue"), legend = c("true", "fitted"), lty = 1, bty = "n", lwd = 4)
```




```
# compute MSE's for d from 1 to 10
degree <- 1:10
cv.mse3 <- sapply(degree, function(d) cv_mse(sim_data3, d))

# best degree
best3 <- degree[which.min(cv.mse3)]
best3
```

Sim data 3

```
## [1] 7
```

```
plot(degree, cv.mse3)
```




```
# fit model using best degree
ols3 <- lm(y ~ poly(x, degree = best3, raw = T), data = sim_data3)

# predicted values
yhat3 <- predict(ols3)

# plot against data and true CEF
plot.new()
plot(y ~ x, sim_data3, pch= 20)
curve(2.83*sin(pi/2*x) + e, col = "red", from = -4, to = 4, add = T, lwd = 4)

## Warning in 2.83 * sin(pi/2 * x) + e: longer object length is not a multiple of
## shorter object length

lines(yhat3[order(x)] ~ sort(x), data= sim_data3, col = "blue", lwd = 4)
legend("topleft", col = c("red", "blue"), legend = c("true", "fitted"), lty = 1, bty = "n", lwd = 4)
```


Sim data 4

```
# compute MSE's for d from 1 to 10
degree <- 1:10
cv.mse4 <- sapply(degree, function(d) cv_mse(sim_data4, d))

# best degree
best4 <- degree[which.min(cv.mse4)]
best4
```

```
## [1] 1
```

```
plot(degree, cv.mse4)
```



```

# fit model using best degree
ols4 <- lm(y ~ poly(x, degree = best4, raw = T), data = sim_data4)


# predicted values
yhat4 <- predict(ols4)

# plot against data and true CEF
plot.new()
plot(y ~ x, sim_data4, pch= 20)
curve(4*sin(3*pi * x) * ifelse(x > 0, 1,0) + e, col = "red", from = -4, to = 4, add = T, lwd = 4)

## Warning in 4 * sin(3 * pi * x) * ifelse(x > 0, 1, 0) + e: longer object length
## is not a multiple of shorter object length

lines(yhat4[order(x)] ~ sort(x), data= sim_data4, col = "blue", lwd = 4)
legend("topleft", col = c("red", "blue"), legend = c("true", "fitted"), lty = 1, bty = "n", lwd = 4)

```



3.2

```

set.seed(42)
n = 1000
x = runif(n, -4, 4)
e = rnorm(n)

#?ifelse

# simulate DGP1
y <- -2*ifelse(x < -3, 1,0) + 2.55 * ifelse(x > -2, 1,0) - 2*ifelse(x > 0, 1,0) +
  4 * ifelse(x > 2, 1,0) - ifelse(x > 3, 1,0) + e
sim_data1 <- data.frame(x, y)

# simulate DGP2

```

```

y <- 6 + 0.4 * x - 0.36 * x^2 + 0.006 * x^3 + e
sim_data2 <- data.frame(x, y)

# simulate DGP3
y <- 2.83*sin(pi/2*x) + e
sim_data3 <- data.frame(x, y)

# simulate DGP4
y <- 4*sin(3*pi *x) * ifelse(x > 0, 1,0) + e
sim_data4 <- data.frame(x, y)

```


DGP

```

# plot size
options(repr.plot.width = 10, repr.plot.height = 10)

# compute MSE's for d from 1 to 10
degree <- 1:10
cv.mse1 <- sapply(degree, function(d) cv_mse(sim_data1, d))
plot(degree, cv.mse1)

```



Sim data 1

```

# best degree
best1 <- degree[which.min(cv.mse1)]
best1

## [1] 10

# fit model using best degree
ols1 <- lm(y ~ poly(x, degree = best1, raw = T), data = sim_data1)

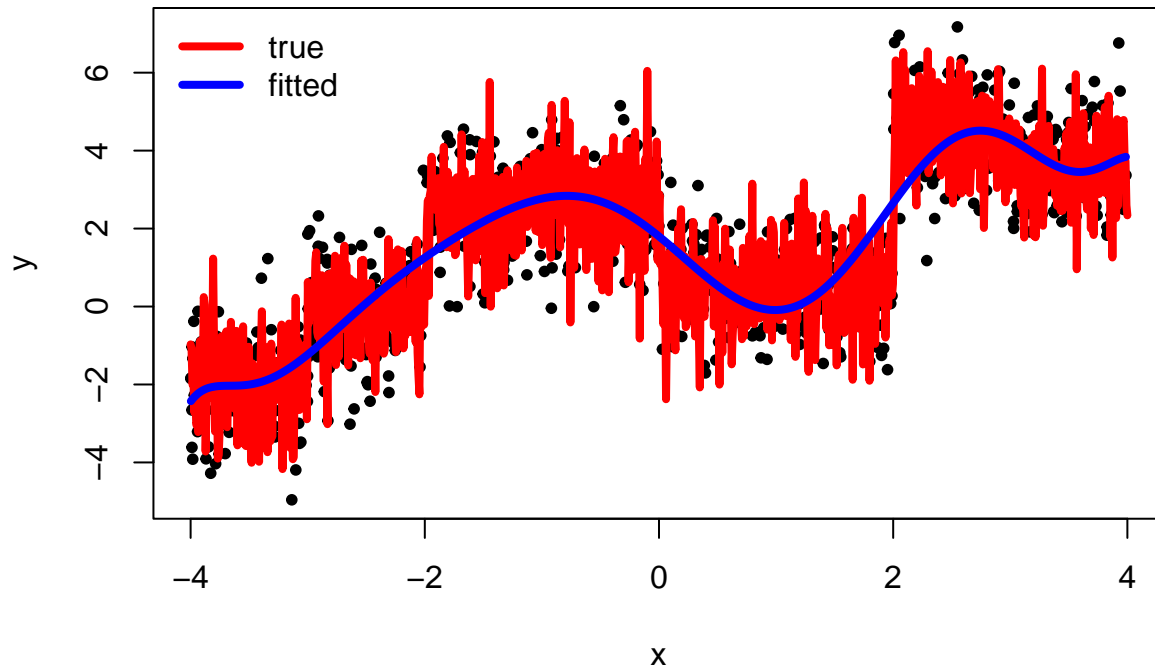
# predicted values
yhat1 <- predict(ols1)

```

```

# plot against data and true CEF
plot.new()
plot(y ~ x, sim_data1, pch= 20)
curve( -2*ifelse(x < -3, 1,0) + 2.55 * ifelse(x > -2, 1,0) - 2*ifelse(x > 0, 1,0)+
      4 * ifelse(x > 2, 1,0) - ifelse(x > 3, 1,0) + e, col = "red", from = -4, to = 4, n =n, add = T)
lines(yhat1[order(x)] ~ sort(x), data= sim_data1, col = "blue", lwd = 4)
legend("topleft", col = c("red", "blue"), legend = c("true", "fitted"), lty = 1, bty = "n", lwd = 4)

```



```

# compute MSE's for d from 1 to 10
degree <- 1:10
cv.mse2 <- sapply(degree, function(d) cv_mse(sim_data2, d))


# best degree
best2 <- degree[which.min(cv.mse2)]
best2

```

Sim data 2

```
## [1] 4
```

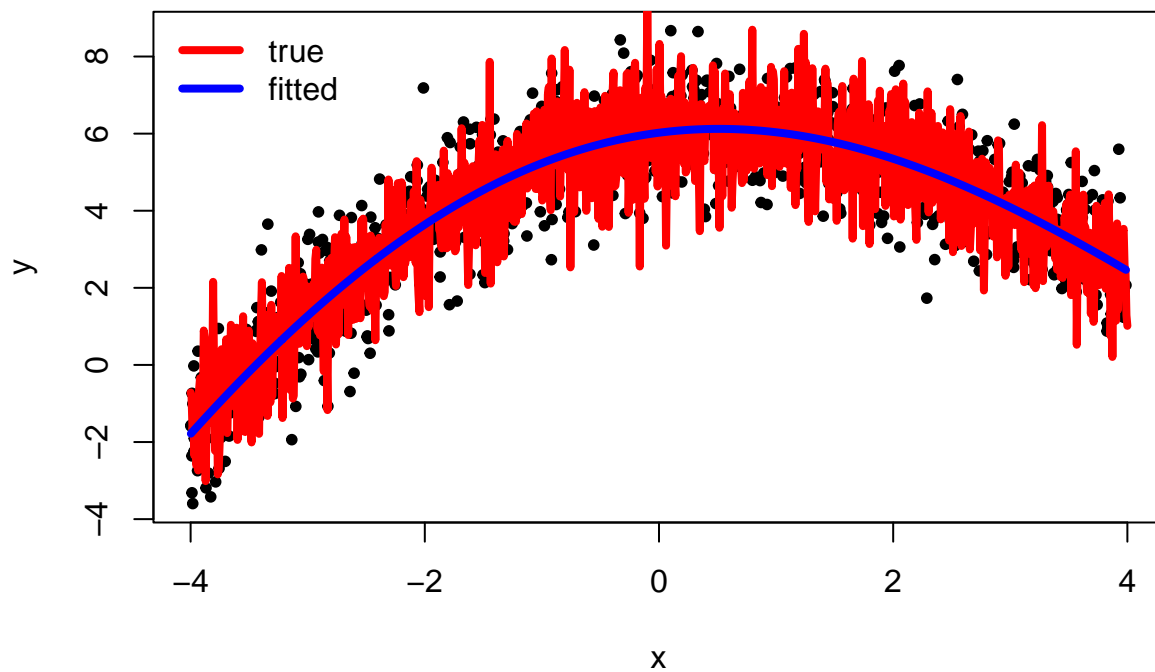
```
plot(degree,cv.mse2)
```



```
# fit model using best degree
ols2 <- lm(y ~ poly(x, degree = best2, raw = T), data = sim_data2)

# predicted values
yhat2 <- predict(ols2)
?curve

# plot against data and true CEF
plot.new()
plot(y ~ x, sim_data2, pch= 20)
curve( 6 + 0.4 * x - 0.36 * x^2 + 0.006 * x^3 + e, col = "red", from = -4, to = 4, n =n, add = T, lwd =
lines(yhat2[order(x)] ~ sort(x), data= sim_data2, col = "blue", lwd = 4)
legend("topleft", col = c("red", "blue"), legend = c("true", "fitted"), lty = 1, bty = "n", lwd = 4)
```




```
# compute MSE's for d from 1 to 10
degree <- 1:10
cv.mse3 <- sapply(degree, function(d) cv_mse(sim_data3, d))

# best degree
best3 <- degree[which.min(cv.mse3)]
best3
```

Sim data 3

```
## [1] 9
```


```
plot(degree, cv.mse3)
```



```
# fit model using best degree
ols3 <- lm(y ~ poly(x, degree = best3, raw = T), data = sim_data3)

# predicted values
yhat3 <- predict(ols3)

# plot against data and true CEF
plot.new()
plot(y ~ x, sim_data3, pch= 20)
curve(2.83*sin(pi/2*x) + e, col = "red", from = -4, to = 4, n=n, add = T, lwd = 4)
lines(yhat3[order(x)] ~ sort(x), data= sim_data3, col = "blue", lwd = 4)
legend("topleft", col = c("red", "blue"), legend = c("true", "fitted"), lty = 1, bty = "n", lwd = 4)
```




```
# compute MSE's for d from 1 to 10
degree <- 1:10
cv.mse4 <- sapply(degree, function(d) cv_mse(sim_data4, d))

# best degree
best4 <- degree[which.min(cv.mse4)]
best4
```

Sim data 4

```
## [1] 8
```


```
plot(degree, cv.mse4)
```

```
# fit model using best degree
ols4 <- lm(y ~ poly(x, degree = best4, raw = T), data = sim_data4)

# predicted values
yhat4 <- predict(ols4)

# plot against data and true CEF
plot.new()
plot(y ~ x, sim_data4, pch= 20)
curve(4*sin(3*pi *x) * ifelse(x > 0, 1,0) + e, col = "red", n=n, from = -4, to = 4, add = T, lwd = 4)
lines(yhat4[order(x)] ~ sort(x), data= sim_data4, col = "blue", lwd = 4)
legend("topleft", col = c("red", "blue"), legend = c("true", "fitted"), lty = 1, bty = "n", lwd = 4)
```



The results are different because more data can better help us fit the models.

Q4

4.1

```
tecator <- read.csv('tecator.csv')
ols_tecator <- lm(fat ~ ., data = tecator)
mean((predict(ols_tecator) - tecator$fat)^2)
```

```
## [1] 0.7898249
```

This possibly is not a good estimate of out-of-sample performance because there are too many predictors.

4.2

```
cv_mse1 <- function(data, k = 5){
  set.seed(59)
  # create folds randomly
  n <- nrow(data)
  folds <- sample(rep(1:k, length = n))

  # create vector to store results
  mse <- rep(NA, k)
  for(j in 1:k){

    # train model on all folds except j
    train <- folds != j
    ols <- lm(fat ~ ., data = data[train, ])

    # compute MSE on fold j (not used for training)
    yhat <- predict(ols, newdata = data[!train, ])
    mse[j] <- mean((data$fat[!train] - yhat)^2)
```

```

}
  # compute average mse
  mse.cv <- mean(mse)
  return(mse.cv)
}
cv_mse1(tecator)

```

```
## [1] 13.60371
```

The MSE is different from the OLS model, and interestingly larger. Possibly due to the variation in predictors.

4.3

```

# load package
library(glmnet)

## Loading required package: Matrix
## Loaded glmnet 4.1-3


# plot size
options(repr.plot.width = 10, repr.plot.height = 10)

# scipen
options(scipen = 99)

# create model matrix
X <- model.matrix(fat ~ . -1, data = tecator)
y <- tecator$fat

# alpha = 1 (Lasso)
lasso <- glmnet(X, y, alpha = 1)
plot(lasso, xvar = "lambda", label = T)

```



```
cv.lasso <- cv.glmnet(X, y, type = "mse", nfolds = 5, nlambda = 100, lambda = seq(0, 1, by = 0.01), alpha = 1)
```

```
cv.lasso
```

```
##
```

```
## Call: cv.glmnet(x = X, y = y, lambda = seq(0, 1, by = 0.01), type.measure = "mse", nfolds = 5,
```

```
##
```

```
## Measure: Mean-Squared Error
```


```
##
```

```
##      Lambda Index Measure      SE Nonzero
```

```
## min      0   101   9.216 0.9126      100
```

```
## 1se      0   101   9.216 0.9126      100
```

```
plot(cv.lasso)
```



```
coef(cv.lasso, s = "lambda.min")
```

```
## 101 x 1 sparse Matrix of class "dgCMatrix"
##                               s1
## (Intercept) 12.33362263711
## X1          26.25016782401
## X2           0.00132744517
## X3           0.00128815092
## X4          -17.14430935904
## X5           24.31156909980
## X6           22.58445823756
## X7           20.97362218397
## X8           20.68858272735
## X9           15.94540874881
## X10          -49.99717113627
## X11          -62.24684034308
## X12          -26.70717792041
## X13          220.84369109788
## X14           -1.14726566823
## X15          -55.65582735566
## X16          -42.25910340554
## X17          -36.23956649454
## X18          -33.22480557033
## X19          -27.90113824952
## X20          -25.14916900822
## X21          -20.87778344594
## X22          -15.80940554138
## X23           -8.59676177743
## X24           1.93375256439
## X25           10.21005429332
## X26           12.18680783500
```


## X27	8.52427175847
## X28	2.25622547849
## X29	-2.05424550914
## X30	-8.82270298516
## X31	-17.27352129986
## X32	-22.41790490740
## X33	-21.05215769630
## X34	-11.45502422799
## X35	-0.06353441172
## X36	9.43569470116
## X37	15.23284643548
## X38	16.60688789648
## X39	10.64857935571
## X40	8.74178860766
## X41	141.38501824871
## X42	0.00013796077
## X43	0.00024574255
## X44	0.00006437054
## X45	-0.00029220556
## X46	-0.00064655588
## X47	-0.00076993003
## X48	-0.00063047886
## X49	0.00023394090
## X50	-31.41070306130
## X51	-62.96802980668
## X52	-78.97039204924
## X53	-8.54225641510
## X54	28.66814057769
## X55	-14.23089446040
## X56	72.74811430234
## X57	4.48529136986
## X58	-2.39057970297
## X59	77.47038921163
## X60	-2.08781949092
## X61	-4.81420422516
## X62	-5.50703021308
## X63	-13.14311379365
## X64	11.49899429745
## X65	-0.00057320931
## X66	-0.00051863842
## X67	-0.00020465717
## X68	-0.00004878772
## X69	-0.00027621292
## X70	-0.00052362621
## X71	-33.01518029121
## X72	-9.04054325787
## X73	-12.54729474338
## X74	-17.29190283596
## X75	-12.33370428880
## X76	0.00016587262
## X77	0.00013818700
## X78	0.00011241482
## X79	0.00009267081
## X80	0.00016158623

```
## X81      0.00006911783
## X82      0.00004084222
## X83     -0.00004035699
## X84     -0.00011224087
## X85     -0.00005391842
## X86     -0.00003013422
## X87    -42.61359398964
## X88      7.27104681312
## X89      7.82257805523
## X90      6.11520300338
## X91      4.43647894288
## X92      5.05177460978
## X93      4.71051086688
## X94      4.83322033403
## X95      5.39841745342
## X96      5.33607589317
## X97      6.42067648638
## X98      7.13753041402
## X99      7.86283199351
## X100     3.77305303478
```

With 5-fold cross validation, the MSE of the best Lasso model where lambda ranges from 0 to 1 is 13.60371 as calculated in b, with lambda = 0.

4.4


```
# alpha = 0 (Ridge)
ridge <- glmnet(X, y, alpha = 0)
plot(ridge, xvar = "lambda", label = T)
```



```
cv.ridge <- cv.glmnet(X, y, type = "mse", nfolds = 5, nlambda = 100, lambda = seq(0, 1, by = 0.01), alpha = 0)
cv.ridge
```

```
##
## Call: cv.glmnet(x = X, y = y, lambda = seq(0, 1, by = 0.01), type.measure = "mse", nfolds = 5,
##
## Measure: Mean-Squared Error
##
##      Lambda Index Measure      SE Nonzero
## min  0.01   100   10.67 1.541     100
## 1se  0.05    96   12.20 1.329     100
```

```
plot(cv.ridge)
```



```
coef(cv.ridge, s = "lambda.min")
```

```
## 101 x 1 sparse Matrix of class "dgCMatrix"
##              s1
## (Intercept) 19.9432711
## X1          40.3608315
## X2          26.6506769
## X3          17.2794168
## X4          10.2135713
## X5           4.7961857
## X6           0.5793649
## X7          -2.7730942
## X8          -5.4241091
## X9          -7.6768084
## X10         -9.4637154
## X11        -10.8909147
```


## X12	-12.0220722
## X13	-12.9041983
## X14	-13.5424727
## X15	-13.9531493
## X16	-14.1055597
## X17	-14.0962039
## X18	-13.8721487
## X19	-13.3610740
## X20	-12.6470420
## X21	-11.6595023
## X22	-10.5373811
## X23	-9.3939194
## X24	-8.2630803
## X25	-7.2364570
## X26	-6.3936161
## X27	-5.5869877
## X28	-4.5390767
## X29	-3.0011738
## X30	-1.0367287
## X31	1.1677970
## X32	3.4789956
## X33	5.7862596
## X34	8.0857936
## X35	10.3171491
## X36	12.5573538
## X37	14.8722742
## X38	17.0526000
## X39	18.6293787
## X40	19.3512865
## X41	19.0980732
## X42	17.7689086
## X43	15.1352395
## X44	11.0101268
## X45	5.6711457
## X46	-0.1601710
## X47	-5.5244397
## X48	-9.4037597
## X49	-11.2009597
## X50	-11.2102509
## X51	-9.9961488
## X52	-7.8053611
## X53	-4.9399039
## X54	-1.8388927
## X55	0.9046726
## X56	2.8768261
## X57	3.8718937
## X58	4.0058566
## X59	3.3660410
## X60	2.7695399
## X61	2.0133849
## X62	1.3099630
## X63	0.5270859
## X64	-0.2216513
## X65	-0.9064106

```
## X66      -1.4395741
## X67      -1.7694084
## X68      -2.0741876
## X69      -2.5878357
## X70      -3.1351353
## X71      -3.4096255
## X72      -3.4150404
## X73      -3.5205399
## X74      -3.7221289
## X75      -3.6488169
## X76      -3.2035240
## X77      -2.9054436
## X78      -2.6006382
## X79      -2.2509559
## X80      -1.6749673
## X81      -1.0620827
## X82      -0.4276202
## X83       0.1515666
## X84       0.6108982
## X85       1.0224524
## X86       1.2426008
## X87       1.2919510
## X88       1.2033113
## X89       1.1480918
## X90       1.1696162
## X91       1.3674577
## X92       1.7525997
## X93       2.1516185
## X94       2.4860797
## X95       2.6981535
## X96       2.7548630
## X97       2.6344771
## X98       2.2392776
## X99       1.5030625
## X100      0.4773476
```

With 5-fold cross validation, the MSE of the best ridge model where lambda ranges from 0 to 1 is 13.60371 as calculated in b, with lambda = 0.

Q5

5.1

The ATE is 0. Because D and Y are independent. We can still run Y on D and we only include D, but the effect is not causal.

5.2

```
lasso <- read.csv('lasso.csv')
cv.lasso1 <- cv.glmnet(X, y, type = "mse", alpha = 1)


cv.lasso1

##
## Call:  cv.glmnet(x = X, y = y, type.measure = "mse", alpha = 1)
```

```
##
## Measure: Mean-Squared Error
##
##      Lambda Index Measure      SE Nonzero
## min 0.01073    70    11.37 1.163        22
## 1se 0.02986    59    12.49 1.072        14
plot(cv.lasso1, xvar = "lambda", label = T)

## Warning in plot.window(...): "xvar" is not a graphical parameter
## Warning in plot.window(...): "label" is not a graphical parameter
## Warning in plot.xy(xy, type, ...): "xvar" is not a graphical parameter
## Warning in plot.xy(xy, type, ...): "label" is not a graphical parameter
## Warning in axis(side = side, at = at, labels = labels, ...): "xvar" is not a
## graphical parameter
## Warning in axis(side = side, at = at, labels = labels, ...): "label" is not a
## graphical parameter
## Warning in axis(side = side, at = at, labels = labels, ...): "xvar" is not a
## graphical parameter
## Warning in axis(side = side, at = at, labels = labels, ...): "label" is not a
## graphical parameter
## Warning in box(...): "xvar" is not a graphical parameter
## Warning in box(...): "label" is not a graphical parameter
## Warning in title(...): "xvar" is not a graphical parameter
## Warning in title(...): "label" is not a graphical parameter
```

22 17 14 8 9 8 7 8 12 9 4 4 1 1 1 1 1



```

coef(cv.lasso1, s = "lambda.min")

## 101 x 1 sparse Matrix of class "dgCMatrix"
##                               s1
## (Intercept) 23.326350444
## X1          46.281285546
## X2          .
## X3          .
## X4          .
## X5          .
## X6          .
## X7          .
## X8          .
## X9          .
## X10         .
## X11         .
## X12         .
## X13         .
## X14        -46.391774302
## X15        -24.146945250
## X16        -13.805976718
## X17        -19.081269344
## X18        -12.459335761
## X19         -5.282139712
## X20         -5.498186167
## X21         -1.928623158
## X22         -0.905030785
## X23         -0.513157798
## X24         -2.978587681
## X25         .
## X26         .
## X27         .
## X28         .
## X29         .
## X30         .
## X31         .
## X32         .
## X33         .
## X34         .
## X35         .
## X36         .
## X37         .
## X38         .
## X39         0.002004740
## X40         2.924055010
## X41        139.783700985
## X42         .
## X43         .
## X44         .
## X45         .
## X46         .
## X47         .
## X48         .
## X49        -0.001367663

```

## X50	-0.176612221
## X51	-30.427759331
## X52	-33.872229527
## X53	.
## X54	.
## X55	.
## X56	.
## X57	.
## X58	.
## X59	.
## X60	.
## X61	.
## X62	.
## X63	.
## X64	.
## X65	.
## X66	.
## X67	.
## X68	.
## X69	.
## X70	.
## X71	.
## X72	.
## X73	.
## X74	.
## X75	.
## X76	.
## X77	.
## X78	.
## X79	.
## X80	.
## X81	.
## X82	.
## X83	.
## X84	.
## X85	.
## X86	.
## X87	.
## X88	.
## X89	.
## X90	.
## X91	.
## X92	.
## X93	.
## X94	.
## X95	.
## X96	.
## X97	.
## X98	0.681081919
## X99	3.400822636
## X100	0.912022256

There are 22 variables selected by Lasso, as shown above.

5.3

No, this is not a good estimate of the causal effect because the coefficients are only used for prediction. It is not specific to Lasso.